

CS1102: Lecture 10



Priority Queues and Heaps

Chapter 12: pages 621-640

Outline

- What is the ADT **priority queue**?
- What are the **operations** supported?
- **heap** (**max**-heap, **min**-heap)
 - **heapInsert** ($O(\log N)$)
 - **heapDelete** ($O(\log N)$)
 - **heapRebuild** ($O(\log N)$)
 - **heapify** ($O(N)$) - **heap construction** algorithm
- **heapSort** ($O(N \log N)$)
- Java API: **PriorityQueue**<E>

What is a **Priority Queue**?

A Special form of queue from which **items are removed according to their designated priority** and not the order in which they entered.

Priority Queue - Examples

- A “to-do” list with priorities
 - Scheduling jobs in operating systems (OS)
- Go to Takashimaya (6)
 - Work out in gym (5)
 - Prepare CS1102 lecture slides (1)
 - Go to department tea (4)
 - ...

Priority queue **operations**

- **Create** an empty priority queue
- **Insert** an item with a given key
- **Remove** the item with **maximum** (or minimum) key value
- Determine whether a priority queue **is empty**.

Priority Queue – **Unsorted list** implementation

- ❑ **Insertion**: insert an element to end of a list. Complexity is $O(1)$.
- ❑ **Removing maximum key value**: traverse the list to find the element of maximum key value and remove it. Complexity is $O(n)$.

Priority Queue – **Sorted list** implementation

- **Insertion:** $O(n)$.
- **Removing maximum key value:** $O(1)$.

Priority Queue –

Binary Search Tree implementation

- ❑ **Insertion:** $O(h)$, where h is the height of the BST.
- ❑ **Removing maximum key value:** $O(h)$.
- ❑ The height of a BST is **not** always $O(\log n)$.

Priority Queue – **AVL Tree** implementation

- ❑ AVL Tree is not covered in this course. Just for your information only.
- ❑ An AVL Tree is a self-balancing **height balanced** binary search tree.
- ❑ The height of an AVL tree is $O(\log n)$.
- ❑ **Insertion**: $O(\log n)$.
- ❑ **Removing maximum key value**: $O(\log n)$.

Priority Queue –

Heap implementation (Summary)

- **Insertion**: $O(\log n)$.
- **Removing maximum key value**: $O(\log n)$.
- **Heap construction** $O(n)$.

How about for building an AVL tree?

Ans: $O(n \log n)$

- **Peek** the maximum key value: $O(1)$.

How about for an AVL tree?


Ans: $O(\log n)$

Heap




Heap is the most appropriate data structure for realizing the ADT priority queue.

Definition: **Heap**



- A (binary) **heap** is a **complete binary** tree
 - either is empty,
 - or satisfies the **heap property**:
for every node v , the search key in v is
greater or equal to those in the **children** of
 v . It is called a **max heap**.


Note: The keys on the right subtree could be smaller than the keys on the left subtree




Q: What are the nice properties of complete binary trees?

Q: Is it a heap?

Negative heap examples



not complete




fail heap
property

Compare **heap** with **binary search tree (BST)**

- Both are binary trees.
- Difference
 - Heap maintains **heap property**.
 - It is not a binary search tree
 - BST maintains **BST property**.
 - It is not a heap

Representation of heaps using arrays



Recall

$$\text{left}(i) = 2*i+1$$


$$\text{right}(i) = 2*i+2$$

$$\text{parent}(i) = \text{floor}((i-1)/2)$$

Note: Recall that a heap is a complete binary satisfying **heap property**

Insert an item


Insert 15



0	9
1	8
2	7
3	4
4	3
5	5
6	1
7	2
8	15

To insert an item, we first **append** it to the array. This may **violate** the heap property.


Re-establish heap property




We **re-establish** the heap property by moving the newly inserted item up the tree. This operation is called **bubble up**.

Re-establish heap property

15 is bubbled up again




Re-establish heap property



15 is bubbled up to its final position. The heap is re-established.

Remove the max item




The max item 15 (the root) is replaced by the last item 1.

To remove the **maximum** item (the root item, **12** in this example), we take the **last item** (**1** in this example) in the heap and replace the **root** with it. This may **violate** the heap property.

Q: Why the **last item**?

Re-establish heap property




Key value **1** violates the heap property

bubble down

If new root node violates the heap property, then we **swap** the new root with its child with **bigger key** value (**9** in this example), and repeat the process until no violation on the heap property. This is called **bubble down**.

Re-establish heap property



bubble down again

Key value **1** is now in its **final position**. The heap property is re-established.

Code - using array to implement heap

```
public class Heap {  
    private int MAX_HEAP = 100;  
    private int [] items;  
    private int size;           // size of the heap  
  
    public Heap() {  
        items = new int[MAX_HEAP];  
        size = 0;  
    }  
  
    public boolean heapIsEmpty() {  
        return size == 0;  
    }  
}
```

heapInsert (bubble up) - using iteration

```
public void heapInsert (int newItem) throws HeapException {
    if (size < MAX_HEAP) {
        items[size] = newItem;           // append the new item to the array
        int place = size;                // place is the index of the new item
        int parent = (place - 1)/2;      // find the parent node index
        while ( (parent >= 0) && (items[place] > items[parent]) {
            // heap property violated, need to bubble up

            int temp = items[place];
            items[place] = items[parent];
            items[parent] = temp;

            place = parent;
            parent = (place - 1)/2;
        }
        ++size;
    }
    else
        throw new HeapException("HeapException: Heap full");
}
```


heapDelete – remove the maximum key value

```
public int heapDelete() {  
    int rootItem = 0;  
    if (!heapIsEmpty()) {  
        rootItem = items[0];    // rootItem is set to the max key value  
        items[0] = items[--size]; //replace the root by the last item  
        heapRebuild (0);        // to rebuild the heap – bubble down  
    }  
    return rootItem;           // return the maximum key value  
}
```

heapRebuild (bubble down) – recursive

- the root node may violate the heap property, bubble down to re-establish the heap recursively

```
protected void heapRebuild (int root) {
    int child = 2 * root + 1;    // left child
    if (child < size) {          // there is a left child
        int rightChild = child + 1; // right child

        if ( (rightChild < size) && // there is a right child
            (items[rightChild] > items[child]) )
            child = rightChild;    // choose child with bigger key value

        if ( items[root] < items[child] ) {
            // bubble down - swap root with bigger child
            int temp = items[root];
            items[root] = items[child];
            items[child] = temp;

            heapRebuild(child);    // recursive call
        }
    }
}
```

Running time of **heapDelete**

- How many calls to **heapRebuild**?
- Go down 1 level after each call to heapRebuild.
- number of heapRebuild calls $<$ **height** of the complete binary tree
- The height of a complete binary tree is $\log n$.
Q: **Why?**
- **Worst case** running time is $O(h) = O(\log n)$
- How about the complexity of **heapInsert**?

Heap Construction




To construct a heap from a given set of key values

Display ranked web pages

- Display 10 pages at a time in order of decreasing page rank scores

Sort the rank scores



- ❑ Sort the web pages according to their rank scores
- ❑ Traverse the sorted list

Running times

- ❑ Sorting $O(n \log n)$
 n : total number of pages
- ❑ Traversing $O(k)$
 k : number of pages requested
- ❑ Total running time:
 $O(n \log n) + O(k)$
 $\leq O(n \log n) + O(n)$, since $k < n$
 $= O(n \log n)$

Q: Can we do better? Maybe

Idea

- Build a **heap** of scores
- Remove the top 10 pages at a time


Heap construction

- Recall **heap property**: for every node v , the search key in v is **greater or equal** to those in the children of v
- Build the heap recursively from **bottom up**

Note: Alternatively, we can insert nodes one at a time into a heap. This corresponds to building a heap from the **top downwards**.

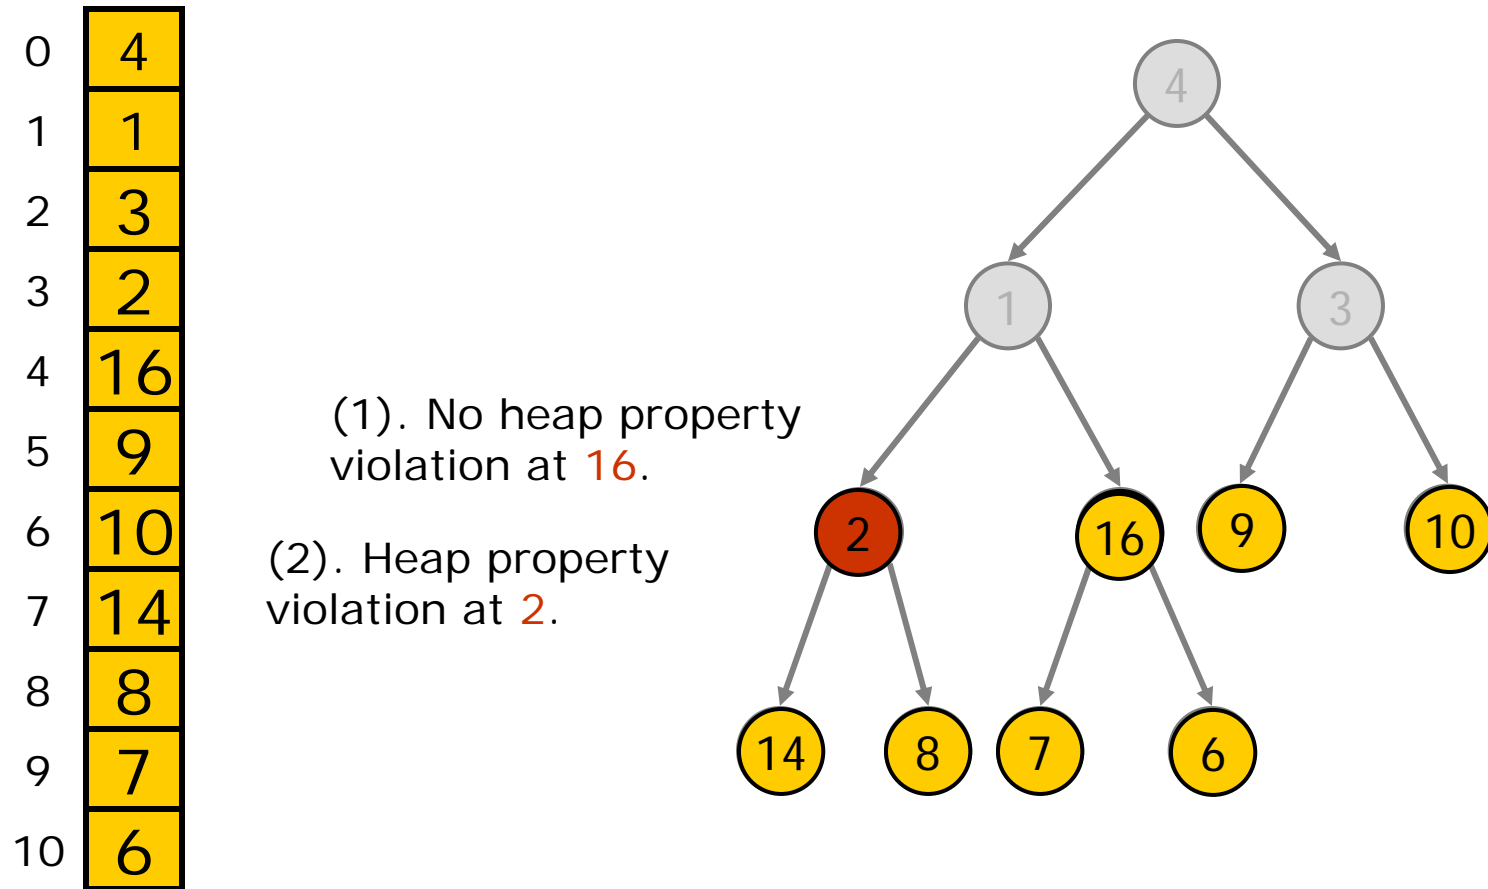
Heap construction example

0	4
1	1
2	3
3	2
4	16
5	9
6	10
7	14
8	8
9	7
10	6



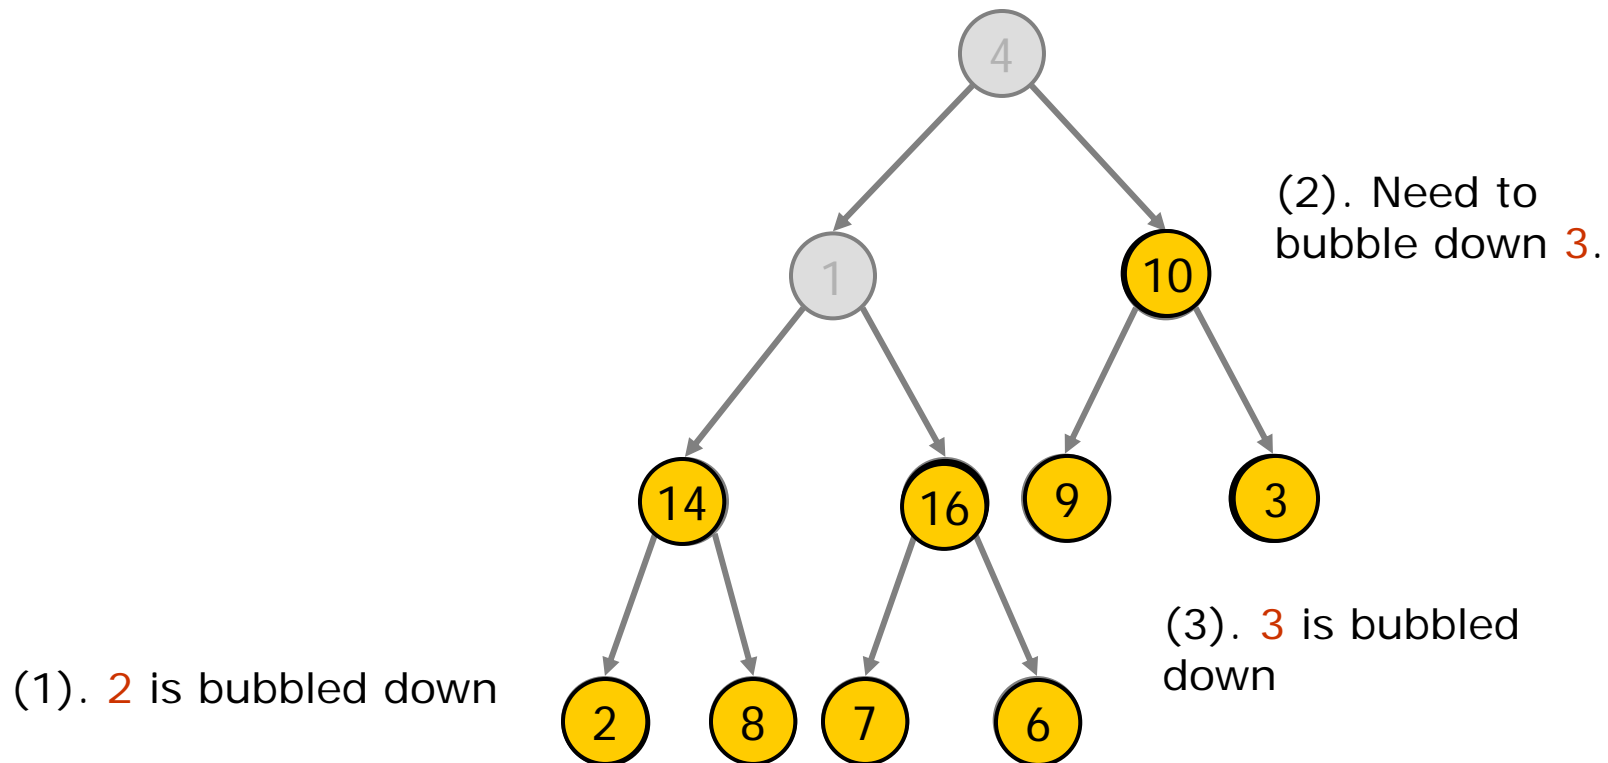
The data in the array represents the complete binary tree on the right. 34

Heap construction example




We start by building a heap by calling the `heapRebuild` method from the **last internal node** back to the **root**. (i.e. **16**, **2**, **3**, **1**, **4** in this example)

Heap construction example



Call **heapRebuild** at 16, then 2, 3, 1, 4 in this example


Heap construction example



Call **heapRebuild** at 16, then 2, 3, 1, 4 in this example

Heap construction example

0	16
1	14
2	10
3	8
4	7
5	9
6	3
7	2
8	4
9	1
10	6



4 is bubbled down. The heap is constructed.

Heapify

- Heap construction algorithm

```
protected void heapify() {  
    for (int i = size/2-1; i >= 0; i--)  
        heapRebuild(i);  
}
```

Note: i starts from the last internal node at $\text{size}/2-1$ back to the root node 0.

Heap construction algorithm's Running time

Rough count: $n/2 * O(\log n) = O(n \log n)$ // call heapRebuild $n/2$ times

More careful count of No. of calls to **heapRebuild**,
level by level from bottom up:

<u>level</u>	<u>No. of calls</u>	<u>Each call requires</u>
-2	$n/2^2$	$O(2)$
-3	$n/2^3$	$O(3)$
-4	$n/2^4$	$O(4)$
...		


Total complexity of the **heap construction algorithm:**

$$2 \times n/2^2 + 3 \times n/2^3 + 4 \times n/2^4 + \dots$$

$$< n(2/2^2 + 3/2^3 + 4/2^4 + \dots)$$

$$< n(3/2) = O(n)$$

Running time: **another derivation**



Count number of edges visited in bubbling down.

number of nodes: $n = 2^h - 1$ // $h = \log n$

Total number of edges visited:

= total no. of edges – no. of edges not visited

= $(n-1) - (h-1)$

= $n-h = O(n)$ **Why?** **Ans.** Because $h = \log n$

Web page ranking again

- Build a heap $O(n)$
 - Retrieve top k pages $O(k \log n)$
Why? *Ans.* heapDelete is $O(\log n)$.
 - Total running time: $O(n) + O(k \log n)$
 - If $k=n$, then $O(n \log n)$ *Q: What meaning?*
 - If $k=20$ (or some other **constant**) , then
$$O(n) + O(20 \log n)$$
$$= O(n)$$
- Retrieve **top K** pages is $O(n)$ using heap!

Heapsort


Uses a heap to sort an array $a[0..n]$ of items.

- ▣ First transform the array into a **heap**, $O(n)$
- ▣ Then execute n steps to turn the heap into a sorted array:
 - In step k , $k = 1..n$:
 - ▣ The array has been partitioned into two regions: the **heap region** $a[0..n-k+1]$ and the **sorted region** $a[n-k+2..n]$
 - ▣ **swap** $a[0]$ with $a[n-k+1]$
 - ▣ **heapRebuild** $a[0..n-k]$.
- ▣ Heapsort worst and average cases are $O(n \log n)$.

Heapsort - Example


Original array

6	3	5	9	2	10
---	---	---	---	---	----



After heap construction


10	9	6	3	2	5
----	---	---	---	---	---




Q: What is complexity of heap construction?

Ans. $O(n)$


Heapsort - Example





Step 1




Before swap



After swap





After heapRebuild



Heapsort - Example

Step 2



Is it in place? **Yes**

Is it stable?

Complexity? **$O(n \log n)$**

After swap

heap				sorted	
2	5	6	3	9	10

After heapRebuild

heap				sorted	
6	5	2	3	9	10

Finally Sorted

heap	sorted				
2	3	5	6	9	10

A Heap of Queues

- ❑ Used when there are only finite number of distinct priority values
 - For example, scheduling of tasks by OS
- ❑ Each entry in the heap is a queue, one queue for each priority
- ❑ To **add** an item to the heap, just enqueue it to the queue of the item's priority
- ❑ To **remove** an item from the heap, dequeue an item from the queue with the **highest priority**

Summary of Heap

- ❑ **Priority queue** is a special form of queue where the item with maximum key value (highest priority) is removed first.
- ❑ It is best implemented by a **Heap**.
- ❑ Heap **Insert** is $O(\log n)$
- ❑ Heap **Delete** is $O(\log n)$
- ❑ **Heap construction** is $O(n)$
- ❑ Peek the **maximum** key value: $O(1)$.
- ❑ Retrieve **Top K** key value: $O(n)$
- ❑ **Heap Sort** is $O(n \log n)$

Q: What are the complexities of the operations if priority queue is implemented by unsorted list, sorted list, binary search tree, or balanced binary tree (e.g. AVL Tree)?

Java API: PriorityQueue<E>

PriorityQueue ()

Creates a PriorityQueue with the default initial capacity (**11**) that orders its elements according to their natural ordering (using Comparable).

PriorityQueue (int initialCapacity, Comparator<? Super E> comparator)

Creates a PriorityQueue with the specified initial capacity that orders its elements according to the specified **comparator**.

PriorityQueue (PriorityQueue<? extends E> c)

Creates a PriorityQueue containing the elements in the specified collection.

Java API: PriorityQueue<E> (cont.)

Java PriorityQueue<E> maintains a **min-heap**.

boolean **offer** (E o)

Inserts the specified element into this priorityqueue.

E **peek**()

Retrieves, but does not remove, the head of this queue, returning null if this queue is empty.

E **poll**()

Retrieves and removes the head of this queue, or null if this queue is empty.

int **size**()

PriorityQueue<E> Example

```
import java.util.*;

public class S {
    public static void main(String[] args) {
        PriorityQueue<Integer> pq = new
            PriorityQueue<Integer>();
        pq.offer(20);
        pq.offer(10);
        while (pq.size() != 0){
            System.out.println( pq.poll() );    // output: 10 20
        }
    }
}
```

Note: Java's PriorityQueue is a **min**-heap.