


Chapter 11: File System Implementation





Chapter 11: File System Implementation

- File-System Structure
- File-System Implementation
- Directory Implementation
- Allocation Methods
- Free-Space Management
- Log-Structured File Systems





Objectives

- To describe the details of implementing local file systems and directory structures
- To describe the implementation of remote file systems
- To discuss block allocation and free-block algorithms and trade-offs





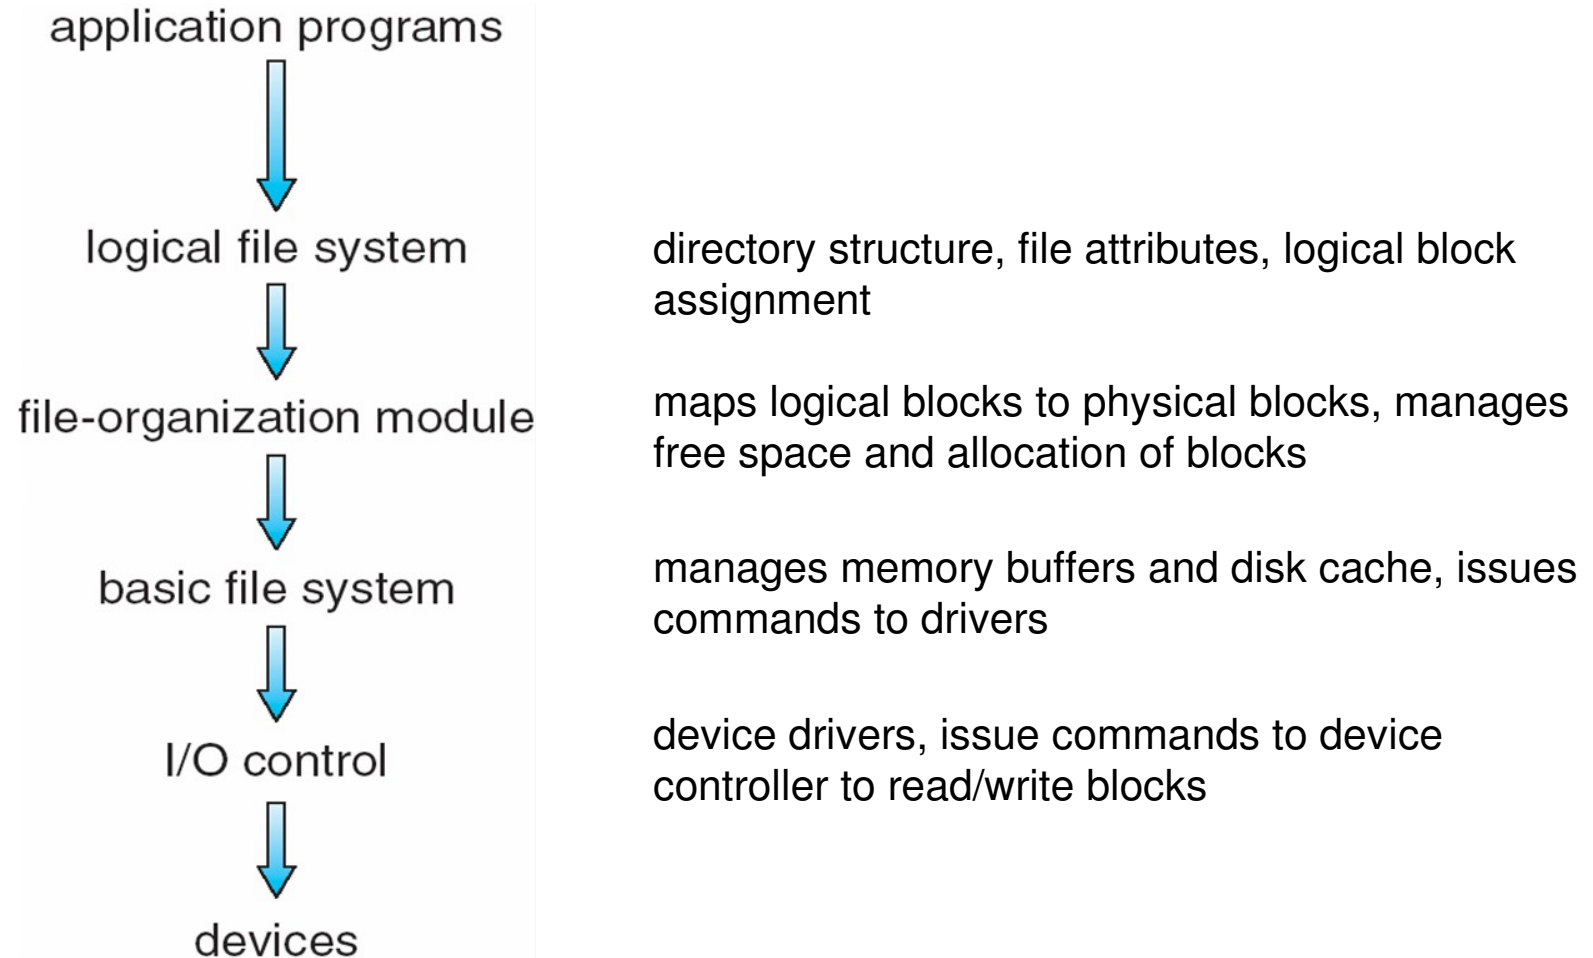
File-System Structure

- File system resides on disks
 - Basic unit is block
 - Random access to blocks
 - I/O transfers are in units of blocks
 - Block maybe spread across multiple sectors
 - Blocks have addresses: drive #, cylinder #, track #, sector #
 - Blocks have logical addresses as well – 0, 1,, 3, ...
- File system organized into layers – see next slide
 - Reduces code redundancy since multiple file systems can share lower levels (e.g. drivers, basic file system)
 - Increases system overhead because of excess calls
- **File control block** – storage structure consisting of information about a file – called inode in UNIX



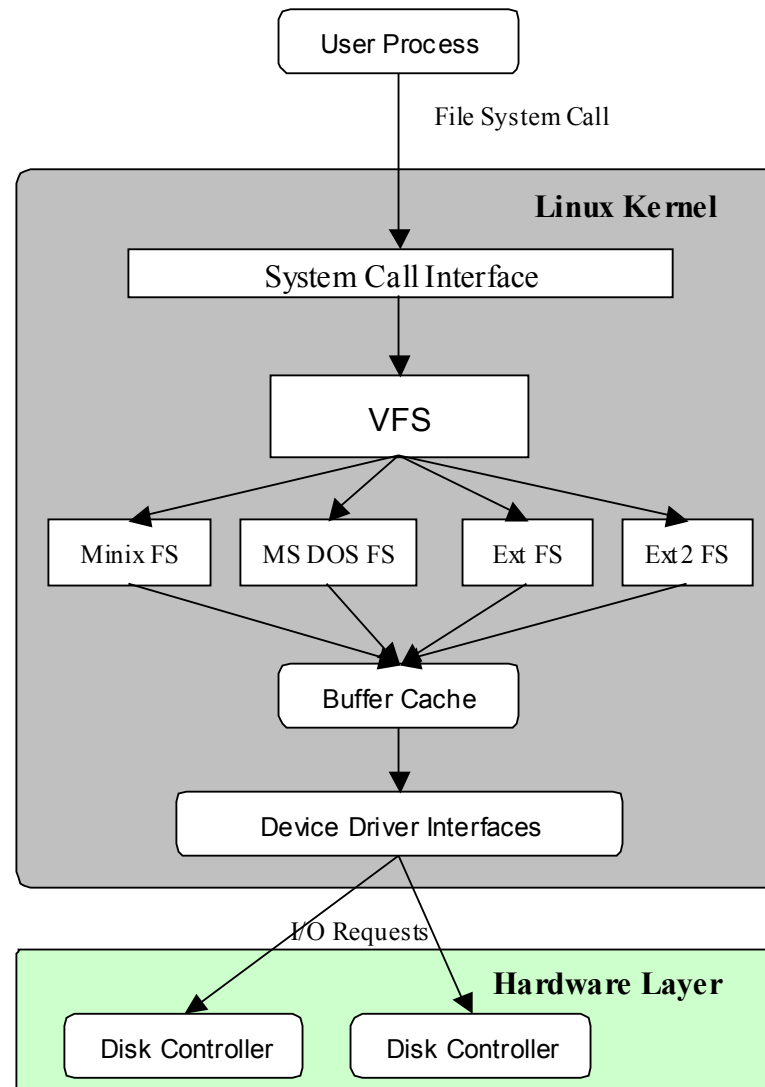


Layered File System





Linux Ext2 File System





A Typical File Control Block

file permissions

file dates (create, access, write)

file owner, group, ACL

file size

file data blocks or pointers to file data blocks





File Systems in Use

- non-Linux UNIX – UNIX File System (UFS), Berkeley Fast File System (FFS),
- Windows – FAT, FAT32, NTFS
- Linux – EXT2, EXT3, and > 40 different systems in general
- All support ISO-9660 (CD-ROMs) and DVD and floppy formats.





File System Implementation

- Basic structures include
 - **Boot control block** (one per volume) – data used by system for booting an operating system from that volume, if an OS is loaded there, otherwise empty. Usually first block on volume. Called **boot block** in UFS, **partition boot sector** in NTFS
 - **Volume control block** (one per volume) – contains number and sizes of blocks in volume, free-block count and list, free FCB count and list; called **superblock** in UFS; in **master file table** in NTFS.
 - **Directory structure** (one per file system) – collection of directories for organizing files – maps file names to FCBs (inodes in UNIX)
 - **File control block** (one per file) – attributes of file: size, location, owner, protections, timestamps, reference counts; called **inode** in UNIX.







In-Memory File Structures

- The OS loads several structures into memory when a file system is mounted. These include
 - **Mount table** – contains information about the volume: access control, disk usage, block sizes, etc.;
 - **Directory structure cache** – a cache of recently accessed directories for faster performance;
 - **System-wide open file table** – copies of FCBs of every open file;
 - **Per-process open file table** – pointers to entries in the system wide open file table for files open by that process
 - Buffers for system blocks for reads/writes from/to disk
- Next slide shows structures used by open() and read() calls





In-Memory File System Structures





File Operations

- To create a file
 - Logical FS is called; allocates new FCB;
 - FCB is filled in with status
 - Data blocks are allocated
 - Address of data blocks filled in FCB
 - Directory in which new name is created is loaded into memory, updated
- To open a file:
 - New slot in per-process open file table created
 - OpSys uses file name to locate FCB (inode) for file;
 - creates new entry in system-wide open file table
 - if file already open, points new entry in system table to FCB in memory
 - if not, loads FCB into memory and points system entry to it
 - process open file table entry points to system table entry





File Operations 2

- To close a file
 - Slot in per-process open file table is deleted
 - Reference count of entry in system-wide open file table is decremented
 - If reference count is zero, metadata is written back to disk and entry is deleted





Partitions and Mounting

- A disk can be divided into multiple partitions
- Each partition can be either raw or have a file system
- Raw partitions are used by databases, for swap space,
- Boot information is stored usually in first few blocks of partition, with its own format, since OpSys is not yet loaded when it is read; the boot loader is stored there, read into memory and executed. The boot loader "knows" file system structure and can read OpSys into memory
- Root partition is the one containing the OpSys. It is mounted at boot time.
- Mounting verifies that device contains valid file system and updates system mount table with info about type of file system, mount point, access attributes, etc





Virtual File Systems

- Virtual File Systems (VFS) provide an object-oriented way of implementing file systems.
- VFS allows the same system call interface (the API) to be used for different types of file systems.
- The API is to the VFS interface, rather than any specific type of file system.





Schematic View of Virtual File System

